

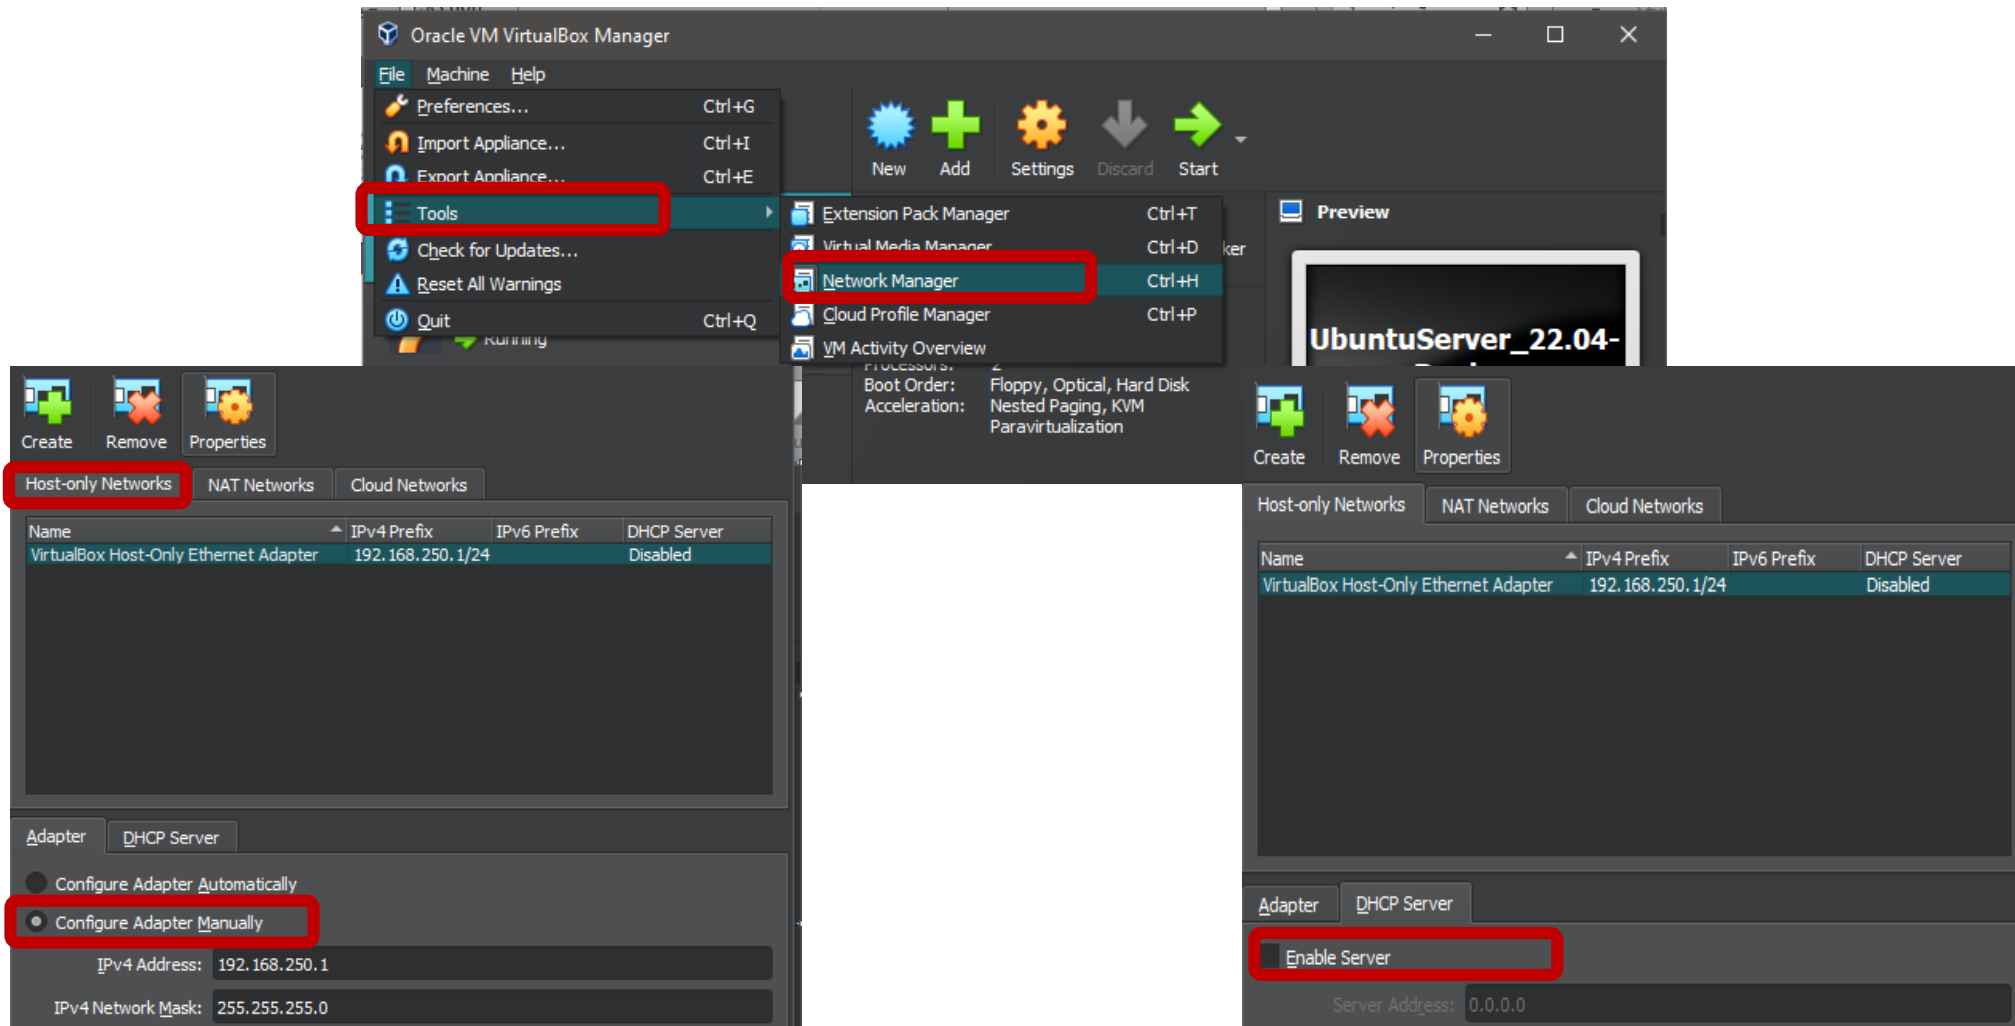
Build Containers with unshare

AU: 2024 - 2025

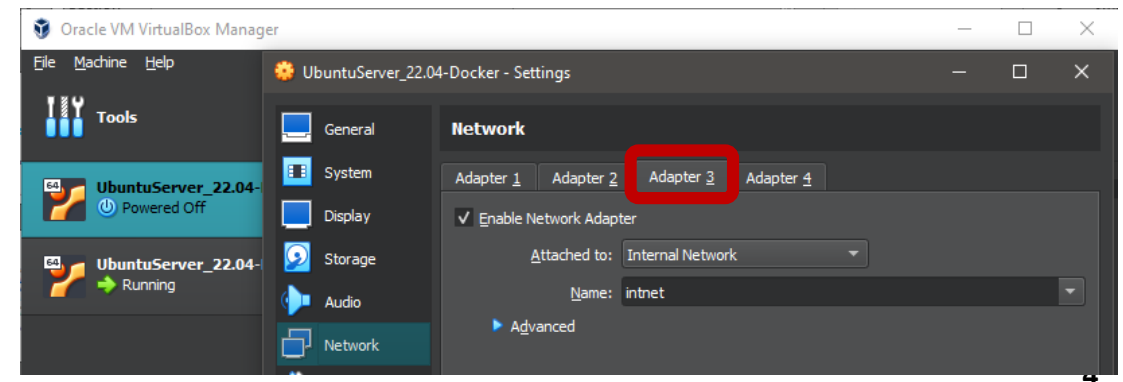
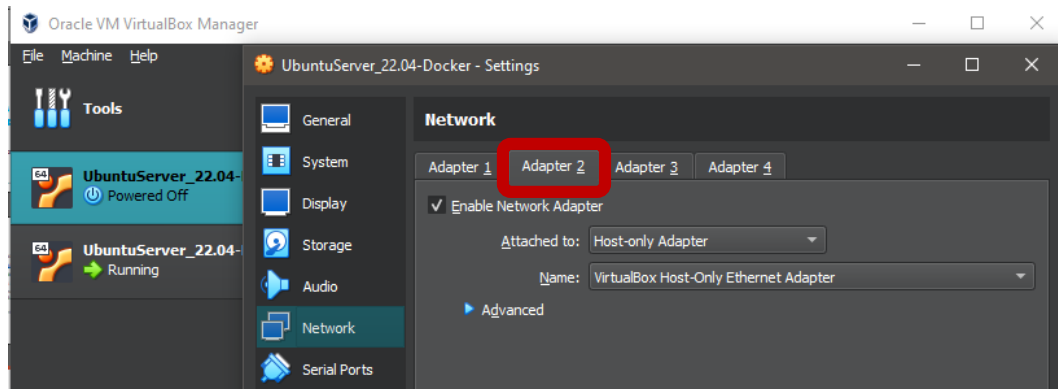
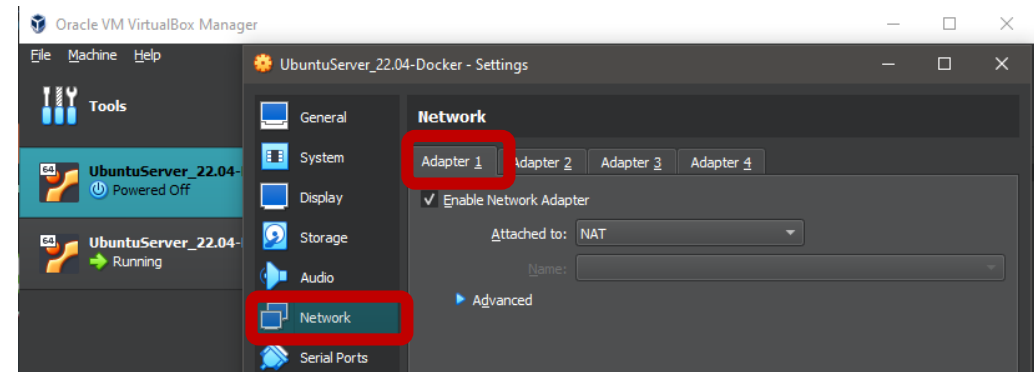
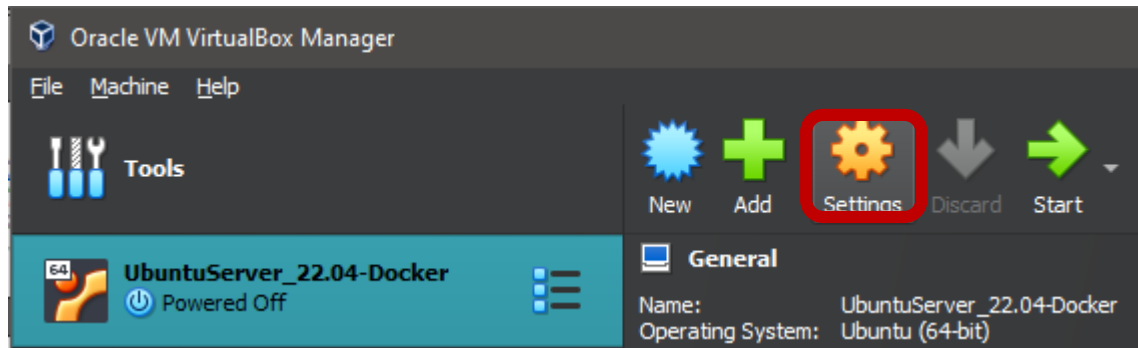
Prerequisites

- Install Linux Ubuntu Server 22.04 or download VM Images:
<https://www.linuxvmimages.com/images/ubuntu-server-2204/>
- **User/Password: ubuntu/ubuntu**
- Network connections of the Linux VM:
 - Network Adapter 1 (enp0s3): NAT (Internet connection)
 - Network Adapter 2 (enp0s8): Host Only Adapter (connection with physical Machine)
 - Network Adapter 3 (enp0s8): Internal Network(connection with VMs)
- Putty installed on physical machine.

Vbox Configuration: Add Host Only Network without DHCP



VM Network Configuration



VM IP Configuration

- Verify the network interfaces (enp0s3,enp0s8 and enp0s9) with: **\$ ip a**
- Edit netplan configuration file: **\$sudo nano /etc/netplan/00-installer-config.yaml**

network:

ethernets:

enp0s3:

dhcp4: true

enp0s8:

dhcp4: false

addresses:

- 192.168.250.10/24

enp0s9:

dhcp4: false

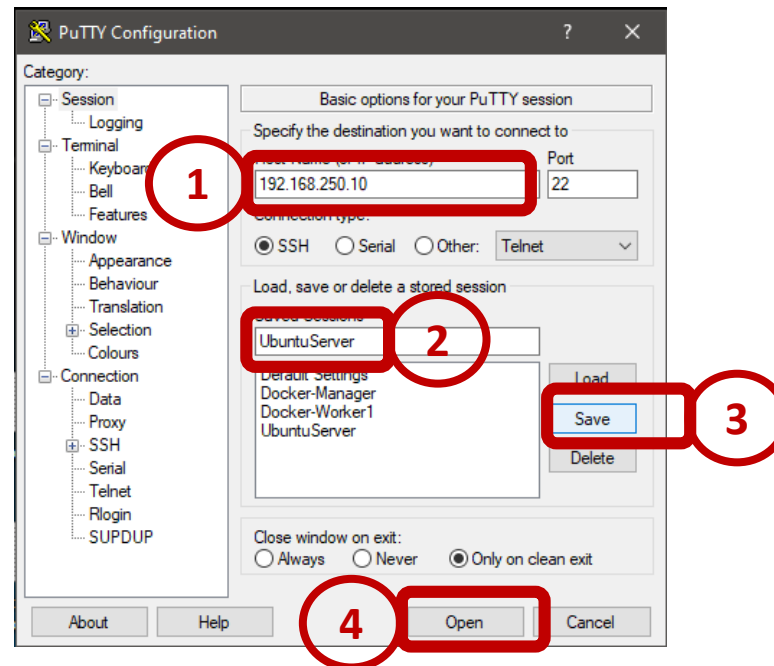
addresses:

- 172.31.0.1/24

version: 2

Connectivity Host - VM

- Apply the new network config: `$sudo netplan apply`
- Verify the VM IP configuration: `$ ip a`
- On physical host:
 - Test ping with the VM, run putty and open an ssh connection:



Creating namespaces with unshare

- Namespaces: Restricted views of systems like the process tree, network interfaces, mounts, ...

- Usage:

unshare [options] [<program> [<argument>...]]

- Options:

-f, --fork: To run the program in the modified namespaces.
unshare will:

- 1) create namespaces,
- 2) fork(),
- 3) Then, exec()

Without -f: Don't fork(): create namespaces and exec() the program directly.

The UTS Namespace

- **Unix Time Sharing (UTS) Namespace:** store the system hostname.

\$ sudo unshare -u bash

❖ The `-u` (`--uts`) flag creates a new UTS Namespace

- Use **pstree** to verify the parent Process (compare with **\$sudo unshare -uf bash**)

#hostname sandbox

❖ To change the hostname in the new UTS name space.

#hostname

❖ Verify the new hostname

- Switch in the native namespaces and check:

#hostname

The UTS Namespace

- **Unix Time Sharing (UTS) Namespace:** store the system hostname.

\$ sudo unshare -u bash

❖ The `-u` (`--uts`) flag creates a new UTS Namespace.

- Use **pstree** to verify the parent Process (compare with **\$sudo unshare -uf bash**)

#hostname sandbox

❖ To change the hostname in the new UTS name space.

#hostname

❖ Verify the new hostname

- Switch in the native namespaces and check:

#hostname

User Namespace

```
$unshare -u bash
```

unshare:unshare failed: Operation not permitted

→ New User namespace: the USER namespace of the created UTS namespace.

```
$unshare -Uu bash
```

- The flag **-U**, **--user**: Create a new USER namespace.
- The flag **-u**, **--uts**: Create a new UTS namespace.

User Namespace

`$unshare -Uu bash` # New user not specified

`$id`

`nobody`

- Display the UID map:

`$cat /proc/self/uid_map` #Empty file: without UID map → nobody in all namespaces (all resources)

`$exit`

`$unshare -Uu --map-user=root bash`

Or: `$unshare -Uu -r bash`: The -r flag and --map-user=root are equivalent.

- The `--map-user=root` or `-r` option: Map current user to root in created name spaces.

`$unshare -Uur bash`

`$whoami`

`root`

`#cat /proc/self/uid_map`

`0 1000` #User with uid 0 is mapped to the user with the UID 1000 (ubuntu)

Verify Namespaces

- The `lsns` command provides information about system namespaces(in native namespaces):

`#lsns`

- With `-p` PID option: Display only the namespaces held by the process with this PID:

`#lsns -p $$`

- Verify that a new **UTS** namespaces is created (compare with the native namespace).

- Reads directly the `/proc` directory:

`#ls -l /proc/Process_ID/ns`

Process_ID can also be: `$$` or `self`

Entering namespaces: nsenter

- Processes may choose to share/join an exited namespaces with **nsenter**.

- In the first terminal:

```
$unshare -uUr bash
#hostname test
#hostname
#echo $$ # PID of container Process
```

- In second terminal, run a bash Sharing the new UTS namespace:

```
$hostname
$sudo nsenter --uts=/proc/ContainerProcess_ID/ns/uts bash
```

```
#Synthax: nsenter --uts=MountNamespaceName Command
# or      : nsenter -u -t ContainerProcess_ID Command
```

```
$hostname
$exit
```

```
sudo nsenter --uts=/proc/ContainerProcess_ID/ns/uts hostname test2
```

- On the first terminal (container): **hostname** # To check the hostname modification

The Mount Namespace: An example

- Simulate disk partition:
 - In the native namespace, create a virtual disk partition:
`$sudo dd if=/dev/zero of=dev1 bs=1024 count=1024`
 - Create the file system (ext2):
`$sudo mkfs -t ext2 dev1`
 - Mount partition, change root owner, create a test file then unmount the partition:
`$sudo mkdir MountPoint`
`$sudo mount dev1 MountPoint`
`$sudo chown ubuntu:ubuntu MountPoint`
`$sudo touch MountPoint/file1`
`$sudo umount MountPoint`

The Mount Namespace: An example

In terminal 1	In terminal 2 (Native namespaces)
<pre> \$unshare -mUf bash # to create a new Mount Namespace for the bash process. ##Synthax: unshare -m (or -mount) command #mount #echo \$\$ # PID of the process # #[Before Running the mount command of terminal 2] #df -kh # list of mounted partitions # ls /home/ubuntu/MountPoint # [After Running the mount command of terminal 2] #df -kh # list of mounted partitions # ls MountPoint #mkdir MountPoint/rep1 #exit </pre>	<pre> df -kh # list of mounted partitions ls MountPoint sudo nsenter -m -t ContainerPID \ mount /home/ubuntu/dev1 \ /home/ubuntu/MountPoint df -kh # list of mounted partitions ls MountPoint </pre>

The PID Namespace

```
$unshare -Umfr -p bash
```

- The flag `-p` (or `-pid`): Create a new PID namespace.
- The `-m` flag: creates a new mount namespace (**why? (1)**)
- The `-f` to fork after creating the new namespaces and starting bash (**why?(2)**)

```
# echo $$
```

```
1
```

```
#ps aux
```

```
# pstree
```

- Why ps report that systemd have the PID 1?
- And the pstree command is referencing the native PID namespace?

The PID Namespace

- The Linux Kernel uses the /proc pseudo filesystem to report running process status.
 - /proc is used by commands like pstree, ps ... to report information about processes

With (unshare -Umfrp bash) /proc is not modified (**The why (1)**).

→ Since we created a new PID namespace, we need to mount a new /proc that matches the new namespace.



The PID Namespace: remount /proc

```
$ unshare -Umfpr bash  
# mount -t proc none /proc  
#ps aux  
# pstree  
#exit
```

- Or directly on command line with **--mount-proc** option:

```
$ unshare -Umfpr --mount-proc bash  
# ps aux
```

The PID Namespace: Why -f option?

Run unshare with (-f)	Run unshare without (-f)
<pre>\$unshare -umUr -p -f sh #mount -t proc none /proc #ls #echo \$\$ 1 #exit</pre> 	<pre>\$unshare -umUr -p sh #mount -t proc none /proc #ls /bin/sh: 2: Cannot fork #echo \$\$ xxx #exit</pre> 

The PID Namespace: Why -f option?

Run unshare with (-f)	Run unshare without (-f)
<pre>#share -umUr -p -f sh #lsns -p \$\$</pre> <p>The sh process is attached to:</p> <ul style="list-style-type: none"> • New PID namespace • New uts namespace • New mount namespace 	<pre>#share -umUr -p sh #lsns -p \$\$</pre> <p>The sh process is attached to a:</p> <ul style="list-style-type: none"> ▪ PID namespace of the parent process (where is the new PID namespace?) ▪ New uts namespace ▪ New mount namespace

Entering in PID namespaces: nsenter

- In a first terminal:

```
$unshare -fUuprm --mount-proc bash
```

```
# tty
```

```
<Terminal name>
```

- In native namespaces: (Identify the PID of bash process)

```
$ ps aux | grep <result of tty command (without /dev)>
```

```
root  xxx 0.0 0.0 20272 3064 pts/5  S+  17:25  0:00 bash
```

```
$pid=xxx
```

```
$sudo nsenter -p -t $pid -m -t $pid unshare -Ufur bash
```

```
#ps aux
```

```
#exit
```

Stop/Re-Start a container

- Stop a running container:

`kill -SIGSTOP $pid`

- Re-start stopped container:

`kill -SIGCONT $pid`

- Run an interactive shell in a container:

`sudo nsenter -a -t $pid sh` #by default the command is bash.

Changing the Filesystem

Example 1: Bash Container

- Create the Directory tree:

```
$TARGET=rootfs
```

```
$mkdir -p ${TARGET}/{dev,proc,bin,lib,lib64}
```

- Determine libraries needed by the bash shell:

```
$ldd /bin/bash
```

```
linux-vdso.so.1 (0x00007ffc3af34000)
```

```
libtinfo.so.6 => /lib/x86_64-linux-gnu/libtinfo.so.6 (0x00007fba7d89000)
```

```
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fba7b60000)
```

```
/lib64/ld-linux-x86-64.so.2 (0x00007fba7f23000)
```

- **linux-vdso.so.1** : A virtual shared object that doesn't have any physical file on the disk.
→ it's a part of the linux kernel.

Changing the Filesystem

Example 1: Bash Container

- Copy files shared object lib

```
$cp -aL /lib/x86_64-linux-gnu/libtinfo.so.6 ${TARGET}/lib/
```

```
$ cp -aL /lib/x86_64-linux-gnu/libc.so.6 ${TARGET}/lib/
```

```
$ cp -aL /lib64/ld-linux-x86-64.so.2 ${TARGET}/lib64/
```

- Copy the binary file

```
$ cp -aL /bin/bash ${TARGET}/bin/
```

- Run container with chroot

```
$ unshare -mipunUrf chroot ${TARGET} /bin/bash
```

```
bash-5.1# echo message
```

```
bash-5.1# pwd
```

```
bash-5.1# ls
```

```
/bin/sh: 1: ls: not found
```

```
bash-5.1# exit
```


Changing the Filesystem

Example 1: Bash Container

- Determine libraries needed by the ls command:

`$ldd /bin/ls`

`linux-vdso.so.1 (0x00007fffccf9c000)`

`libselinux.so.1 => /lib/x86_64-linux-gnu/libselinux.so.1 (0x00007fdbe9896000)`

`libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fdbe966d000)`

`libpcre2-8.so.0 => /lib/x86_64-linux-gnu/libpcre2-8.so.0 (0x00007fdbe95d6000)`

`/lib64/ld-linux-x86-64.so.2 (0x00007fdbe98ed000)`

- Copy files shared object lib used by ls:

`$cp -aL /lib/x86_64-linux-gnu/libselinux.so.1 ${TARGET}/lib/`

`$cp -aL /lib/x86_64-linux-gnu/libpcre2-8.so.0 ${TARGET}/lib/`

- Copy ls binary file:

`$cp -aL /bin/ls ${TARGET}/bin/`

- Re-Run the container:

`$unshare -mipunUrf chroot ${TARGET} /bin/bash`

`bash-5.1# pwd`

`bash-5.1# ls -l`

`bash-5.1# exit`

Changing the Filesystem

Example 2: Debian Container

- Create the root file system:
 - Install debootstrap : A tool that install a Debian-based system into a subdirectory of already-installed system
`$sudo apt install debootstrap`
 - create a new ~/chroot-debian the root directory
`$TARGET="chroot-Debian"`
`$mkdir -p ${TARGET}`
 - With debootstrap, install the debian distribution in TARGET:
`$sudo debootstrap stable ${TARGET} https://deb.debian.org/debian`
 - Verify installed files:
`$ls ${TARGET}`

Changing the Filesystem

Example 2: Debian Container

- Modify the owner/group of the file system:

```
$uid=$(id -u)
```

```
$gid=$(id -g)
```

```
$sudo chown -R $uid:$gid ${TARGET}
```

- Create a Debian Container:

```
$ unshare -mipunUrf chroot ${TARGET} /bin/bash
```

```
    #echo $$
```

```
    #cat /etc/os-release
```

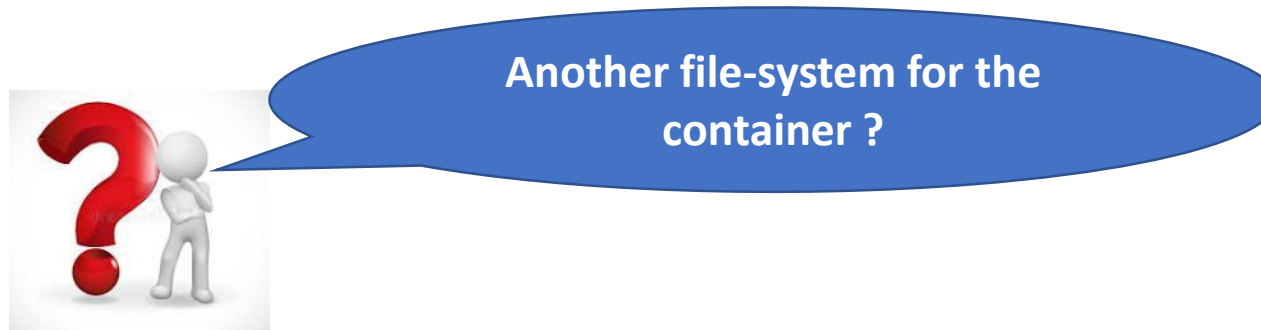
```
    #exit
```

Containers with conventional Linux file-systems

- To run a second bash container:
 - Copy rootfs directory in rootfs2 : `cp rootfs rootfs2`
 - Start the container: `chroot rootfs2 /bin/bash`

☹ **inefficient disk space optimization**: common files branches (exp:/bin) can't be shared by containers.

☹ With (n) instances of a container and a container file system size (s)G, (n*s) G of physical memory would be reserved by running containers.



UnionFS : A File System of Containers

UFS: Union File Systems

chroot

Namespace



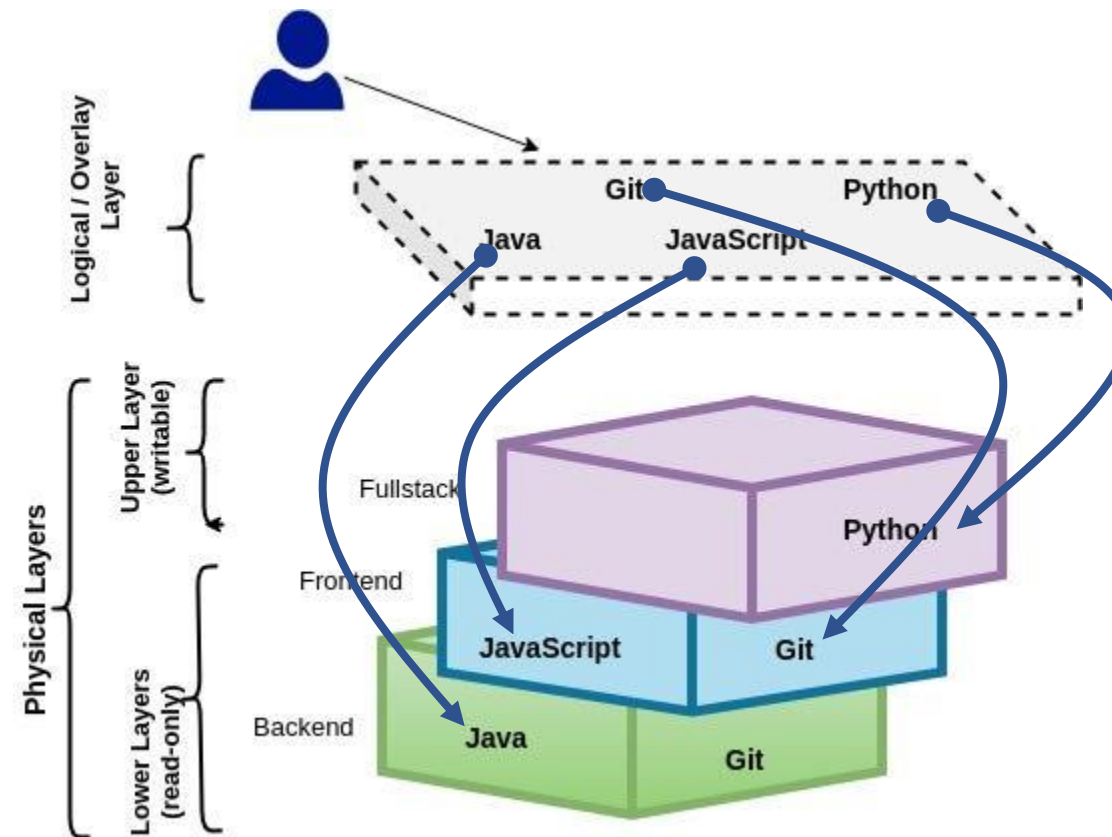
Union File System

- AUFS (advanced multi-layered unification filesystem)
- Overlay FS
- Overlay2 FS
- etc...

Union File System

- **Logical merge** of multiple layers (Layer = branch).
- **Multiple** Read-only **lower layers**, **One** writable **upper layer**.
- Start reading from the **upper layer** than defaults to **lower layers**.
- Copy on Write (CoW) to the upper layer:
A modified file is copied from "lower" to "upper" layer. All the modification will take place in upper layer: the only layer with write access.
- Simulate removal from lower directory through **whiteout** file.

Union File System: A Use Case



Debian Container with UFS

- Create container bundle directory:

```
$TARGET=C1_Debian_Bundle
```

```
$mkdir -p ${TARGET}/{UPPER,WORK,ROOTFS} #UPPER et WORK must be in the same filesystem
```

- Create an OverlayFS (lower layer: rootfs, upperlayer:UPPER):

```
$sudo mount -t overlay \
    -o lowerdir=chroot-Debian,upperdir=${TARGET}/UPPER,workdir=${TARGET}/WORK \
    none ${TARGET}/ROOTFS
```

- Verify the UFS File system

```
$ls ${TARGET}/ROOTFS
```

- Run the Container:

```
$unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash
# echo $$
```

Debian Container with UFS

- Create a second Debian container: Same manip. with (**TARGET=C2_Debian_Bundle**).

- With C1_Debian:

```
# echo Conatiner1 > fileC1
```

```
# rm /etc/os-release
```

- With C2_Debian:

```
# echo Conatiner2 > fileC2
```

```
# cat /etc/os-release
```

- List the Upper layer of C1 container:

```
$ls -l C1_Debian_Bundle/UPPER/
```

```
drwxr-xr-x 2 ubuntu ubuntu 4096 Sep 14 14:03 etc
```

```
-rw-rw-r-- 1 ubuntu ubuntu 11 Sep 14 14:02 fileC1
```

```
$ls -l C1_Debian_Bundle/UPPER/etc
```

```
c----- 2 root root 0, 0 Sep 14 14:03 os-release
```

→ A whiteout file in upper layer (/etc/os-release): Block the visibility of the file.

Network Namespace

- Run two Debian Containers (C1 & C2) on 2 terminals:
 - C1 Debian container:
`TARGET=C1_Debian_Bundle; unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash`
 - C2 Debian container:
`TARGET=C2_Debian_Bundle; unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash`
 - With C1 (or C2):
`# ip a`
1: lo:..... (Only loopback interface)
- Unshare with **-n** option: create an anonymous network namespace.
 - Can be referred by the PID of one process in that namespace.
 - Named network name space are under : `/var/run/netns/`

Network Namespace

- Switch terminal in the native namespaces:
 - The PID(s) of C1 unshare processes (var:pidC1):
`TARGET=C1_Debian_Bundle; pidC1=$(pgrep -f "unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash")`
 - Find the PID of C2 unshare process(var:pidC2 var) :
`TARGET=C2_Debian_Bundle; pidC2=$(pgrep -f "unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash")`
 - Verify pidC1 and pidC2:
`echo -e " - C1 PID: $pidC1 \n - C2 PID: $pidC2"`

Network Namespace

In the native namespaces:

- Create a virtual switch to connect containers (C1 and C2):

```
sudo ip link add br_1 type bridge
```

```
sudo ip link set br_1 up
```

```
sudo ip addr add 10.10.10.1/24 dev br_1
```

- Create pair of virtual ethernet interfaces to connect C1 to switch:

```
sudo ip link add veth_01 netns $pidC1 type veth peer veth_11
```

```
sudo ip link set veth_11 master br_1
```

```
sudo ip link set veth_11 up
```

- Create pair of virtual ethernet interfaces to connect C2 to switch:

```
sudo ip link add veth_02 netns $pidC2 type veth peer veth_12
```

```
sudo ip link set veth_12 master br_1
```

```
sudo ip link set veth_12 up
```

Network Namespace

- **C1 container:**

- `ip address add 10.10.10.2/24 dev veth_01`

- `ip link set veth_01 up`

- To verify the IP configuration: `ip a`

- **C2 container:**

- `ip address add 10.10.10.3/24 dev veth_02`

- `ip link set veth_02 up`

- To verify the IP configuration: `ip a`

Container Networking

- Test connectivity:
- **C1 container (respectively C2):**
 - `ping 10.10.10.1`
 - `ping 10.10.10.3` (respectively `10.10.10.2`)
- **In the native namespaces**, verify that ACCEPT is the default policy in iptables:
 - `sudo iptables -L`
 - `sudo iptables -P FORWARD ACCEPT`

Expose a network port: iptables

- Run/simulate web service with netcat (nc):
`nc [<options>] <host> <port>`
- The nc command: manipulate (read/write) a TCP socket (by default)
- Netcat has two working modes:
 - Listen mode (-l option): Server. If <host> is omitted, nc listens on all addresses
 - Connect mode (the default mode): Client.
- The -k option: When a connection is completed, listen for another one.

Expose a network port: iptables

- Install nc command on container:
 - Exit containers(C1 and C2).
 - Unmount the UFS filesystem (layer0 will be modified):
TARGET=C1_Debian_Bundle; sudo umount \${TARGET}/ROOTFS
TARGET=C2_Debian_Bundle; sudo umount \${TARGET}/ROOTFS
 - Copy nc libraries and binary from VM File system to Layer 0 Container FS :
cd
RootFS=chroot-Debian
cp -afL /lib/x86_64-linux-gnu/libbsd.so.0 \${RootFS}/lib/
cp -afL /lib/x86_64-linux-gnu/libresolv.so.2 \${RootFS}/lib/
cp -afL /lib/x86_64-linux-gnu/libc.so.6 \${RootFS}/lib/
cp -afL /lib/x86_64-linux-gnu/libmd.so.0 \${RootFS}/lib/
cp -aL /usr/bin/nc \${RootFS}/usr/bin/

Expose a network port: iptables

- Remount File Systems/Run container:
- Terminal 1:

```
TARGET=C1_Debian_Bundle
```

```
sudo umount ${TARGET}/ROOTFS
```

```
sudo mount -t overlay -o lowerdir=chroot-Debian,upperdir=${TARGET}/UPPER,workdir=${TARGET}/WORK \
    none ${TARGET}/ROOTFS
```

```
unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash
```

- Terminal2:

```
TARGET=C2_Debian_Bundle
```

```
sudo umount ${TARGET}/ROOTFS
```

```
sudo mount -t overlay -o lowerdir=chroot-Debian,upperdir=${TARGET}/UPPER,workdir=${TARGET}/WORK \
    none ${TARGET}/ROOTFS
```

```
unshare -mipunUrf chroot ${TARGET}/ROOTFS /bin/bash
```

Expose a network port: iptables

In the native namespaces:

```
TARGET=C1_Debian_Bundle; pidC1=$(pgrep -f "unshare -mipunUrf chroot ${TARGET}/ROOTFS  
/bin/bash")
```

```
TARGET=C2_Debian_Bundle; pidC2=$(pgrep -f "unshare -mipunUrf chroot ${TARGET}/ROOTFS  
/bin/bash")
```

- Connect C1 to br_1 switch:

```
sudo ip link add veth_01 netns $pidC1 type veth peer veth_11
```

```
sudo ip link set veth_11 master br_1
```

```
sudo ip link set veth_11 up
```

- Connect C2 to br_1 switch:

```
sudo ip link add veth_02 netns $pidC2 type veth peer veth_12
```

```
sudo ip link set veth_12 master br_1
```

```
sudo ip link set veth_12 up
```

Expose a network port: iptables

- Containers IP configuration:

- ❖ **C1 container:**

- `ip address add 10.10.10.2/24 dev veth_01`

- `ip link set veth_01 up`

- `ip route add default via 10.10.10.1`

- Verify the IP configuration: `ip a`
 - Verify the routing table: `ip route show`

- ❖ **C2 container:**

- `ip address add 10.10.10.3/24 dev veth_02`

- `ip link set veth_02 up`

- `ip route add default via 10.10.10.1`

- Verify the IP configuration: `ip a`
 - Verify the routing table: `ip route show`

Expose a network port: iptables

- Start web service with nc:

- ❖ C1 container:

```
while true;  
do echo -e "HTTP/1.1 200 OK\n\n$(echo 'Debian Container 1' )" \  
| nc -l -k -p 8080 -q 1;  
done
```

- **Ctrl+Z** to stop process

- ❖ C2 container:

```
while true;  
do echo -e "HTTP/1.1 200 OK\n\n$(echo 'Debian Container 2' )" \  
| nc -l -k -p 8090 -q 1;  
done
```

Expose a network port: iptables

- In native namespace(VM Machine): Test web server

```
$curl http://10.10.10.2:8080
```

```
$curl http://10.10.10.3:8090
```

- Expose TCP port:

- C1 container: 8080 – 80

- C2 container: 8090 – 8080

- Configure iptables:

```
sudo iptables -F
```

```
sudo iptables -F -t nat
```

```
sudo iptables -t nat -A PREROUTING -p tcp -d 192.168.250.10 --dport 80 -j DNAT --to-destination 10.10.10.2:8080
```

```
sudo iptables -t nat -A PREROUTING -p tcp -d 192.168.250.10 --dport 8080 -j DNAT --to-destination 10.10.10.3:8090
```

Expose a network port: iptables

- On physical machine:

